Adaptive execution horizon: success versus inference cost (95% Wilson intervals)
PickCube-v1
(test seeds 300000-300199)
PushCube-v1
(fresh seeds 300000-300199)
PullCubeTool-v1
(test seeds 200000-200199)
success rate (success_once)
mean policy inference calls per episode
fixed
random
adaptive_push
adaptive_time
adaptive_dist
fixed(k=8)
fixed(k=4)
fixed(k=2)
adaptive_push
random(hi=8,lo=2,mean=3.76)
random(hi=8,lo=2,mean=2.89)
adaptive_push
adaptive_push
adaptive_push
adaptive_time
adaptive_time
random(hi=8,lo=2,mean=2.25)
fixed(k=2)
adaptive_dist(d=0.3,long=8,short=2)
random(hi=8,lo=2,mean=7.91)
fixed(k=8)
adaptive_dist(d=0.1,long=8,short=2)
adaptive_dist(d=0.4,long=8,short=2)
random(hi=8,lo=2,mean=5.49)
fixed(k=4)
fixed(k=2)