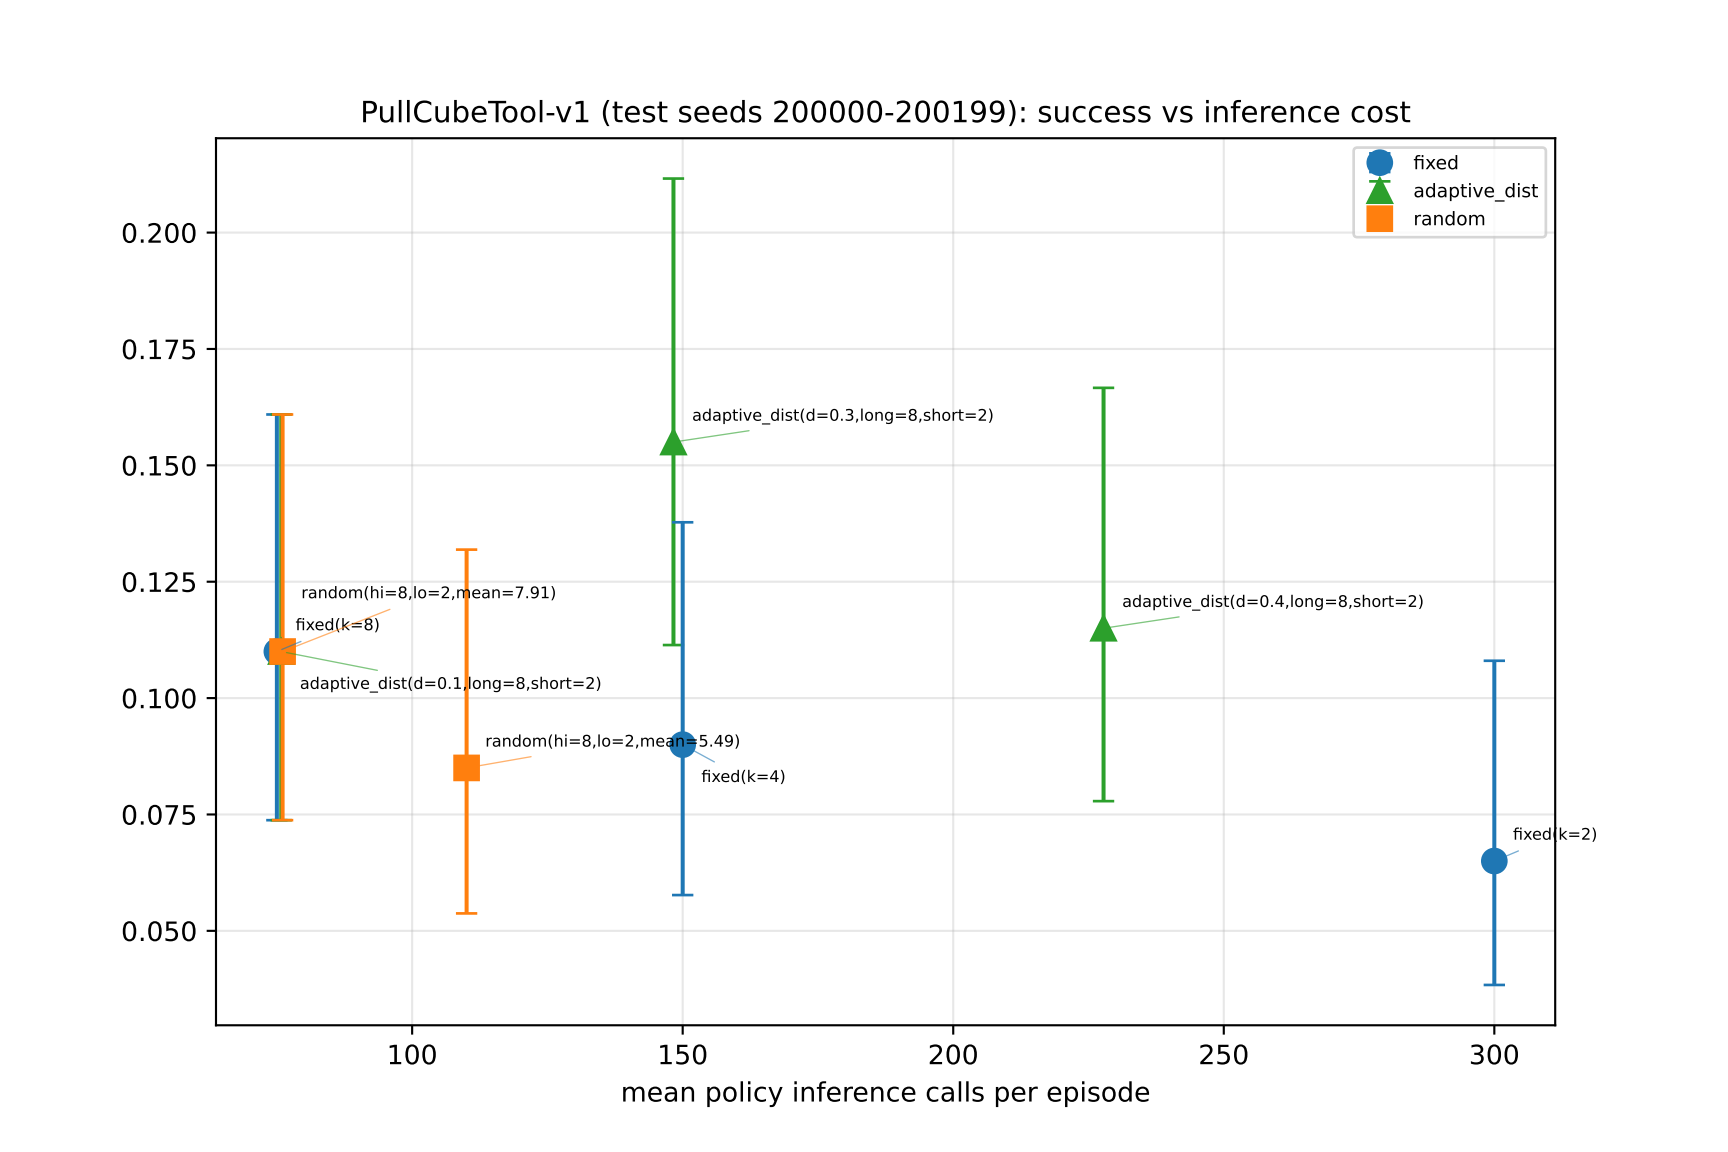
PullCubeTool-v1 (test seeds 200000-200199): success vs inference cost
fixed
adaptive_dist
random
0.200
0.175
0.150
0.125
0.100
0.075
0.050
adaptive_dist(d=0.3,long=8,short=2)
adaptive_dist(d=0.4,long=8,short=2)
random(hi=8,lo=2,mean=7.91)
fixed(k=8)
adaptive_dist(d=0.1,long=8,short=2)
random(hi=8,lo=2,mean=5.49)
fixed(k=4)
fixed(k=2)
100
150
200
250
300
mean policy inference calls per episode