Parity across the three tasks: what the evidence supports, does not establish, or contradicts

	PushCube-v1	PickCube-v1	PullCubeTool-v1
best fixed vs adaptive (co-trained checkpoint) fresh blocks, exact McNemar, n=200 per cell	**supported**	**contradicted**	**not established**
sampler x horizon frontier (co-trained checkpoint) DDPM-20 / DDIM-20 vs DDPM-100: no significant loss, bounded	**supported**	**not established**	**not established**
observation normalization (V1 vs B0) run-level exact permutation test	**supported**	**contradicted**	**supported**
V2a shorter prediction horizon (8) descriptive flag, 3 runs vs 4 (smallest p = 0.057)	**not established**	**suggestive (negative)** -39.9 pp (suggestive)	**suggestive (negative)** -6.3 pp (suggestive)
V2b longer prediction horizon (32) descriptive flag, 3 runs vs 4 (smallest p = 0.057)	**not established**	**not established** +13.6 pp (within)	**suggestive (positive)** +8.0 pp (suggestive)
V3 wider encoder (256) descriptive flag, 3 runs vs 4 (smallest p = 0.057)	**not established**	**not established** +4.0 pp (within)	**suggestive (positive)** +14.7 pp (suggestive)
adaptive replanning on our trained models run-level sign-flip / permutation, 4-6 runs	**contradicted**	**contradicted**	**contradicted**
combination incl. DDIM-5 on trained models pre-specified combination vs default DP, run level	**supported**	**contradicted**	**not established**
co-training vs single-task (multi-run) cut by the user (amendment 7a) - not analysed	**n/a**	**n/a**	**n/a**

supported · not established · contradicted · descriptive only